

# Visualization (Exploring variation)

Peter Ganong and Maggie Shi

January 25, 2026

# Introduction

# Roadmap of lecture

- Introduce three key ideas in data visualization
  - exploration vs. production
  - headline vs. sub-messages
  - iteration
- introduce datasets for this lecture
- Categorical variables
- Continuous variables
  - Exploring typical values
  - Exploring and dealing with unusual values

# What is exploratory data analysis?

Data visualization has two distinct goals

1. **exploration** for you to *learn* as much as possible
2. **production** for you to *teach* someone else what you think the key lessons are

# Exploration vs. production

- When you are in exploration mode, you will look at lots of patterns and your brain filters out the noise
- Production mode is like putting a cone on your dog.
- You are deliberately limiting the reader's field of vision such that they see the key messages from the plot *and avoid too many distractions*

# “A Sunday on La Grande Jatte” by Seurat

# “A Sunday on La Grande Jatte” by Seurat

# **Headline vs submessages**

- **Headline** is what you see first when you look at the painting or you look at the plot.
- **Submessages** are what you see later, on closer inspection
- Often, the most interesting patterns in the data are ones that you don't see right away when you make the very first plot.

# Iterating on plot design

“Make dozens of plots”

Quoctrung Bui, former 30538 guest lecturer and former Harris data viz instructor

# Iterating on plot design

What does he mean?

- The first plot you make will never be the one you should show
- As you are generating graphs for yourself in exploration mode, you will produce many candidates that could end up being used in production mode
- As a rule of thumb, you should try out at least three different plotting concepts (marks)
- Within each concept, you will need to try out several different encodings

# Summary:

1. Decide if you are trying to explore the data or produce a plot for someone else
2. For any given plot, look closely (like Seurat) beyond just the headline
3. Iterate

# Intro to data

# Introduction to data

- Most of our visualization lectures are based on the [University of Washington](#) textbook, but the textbook doesn't have enough material on exploratory data analysis. So we are supplementing with
  - [Data Visualization](#)
  - [Exploratory Data Analysis](#) material in the [R for Data Science](#) textbook (with the code translated to Altair)
- [diamonds](#), [mpg](#) are from “Exploratory Data Analysis”
- [movies](#) (also used last lecture) is from the UW textbook
- [penguins](#) is from “Data Visualization”

# Categorical variables

# Categorical variables: roadmap

- introduce diamonds
- show table
- show bar graph

# introduce dataset diamonds

```
1 import pandas as pd  
2 from plotnine.data import diamonds  
3 diamonds.head()
```

	carat	cut	color	clarity	depth	table	price	x	y
0	0.23	Ideal	E	SI2	61.5	55.0	326	3.95	3.98
1	0.21	Premium	E	SI1	59.8	61.0	326	3.89	3.84
2	0.23	Good	E	VS1	56.9	65.0	327	4.05	4.07
3	0.29	Premium	I	VS2	62.4	58.0	334	4.20	4.23
4	0.31	Good	J	SI2	63.3	58.0	335	4.34	4.35

# introduce dataset diamonds

```
1 diamonds.shape
```

```
(53940, 13)
```

# iamonds data dictionary

Variable	Definition	Values
price	price in USD	\$326–\$18,823
carat	weight of diamond	0.2-5.01
cut	quality of the cut	Fair, Good, Very Good, Premium, Ideal
color	diamond color	D (best) to J (worst)
clarity	measure of how clear diamond is	I1 (worst), SI2, SI1, VS2, VS1, VVS2, VVS1, IF (best)
x	length in mm	0-10.74
y	width in mm	0-58.9
z	depth in mm	0-31.8
depth	$\frac{z}{\text{mean}(x,y)}$	43-79
table	width of top of diamond relative to widest point	43-95

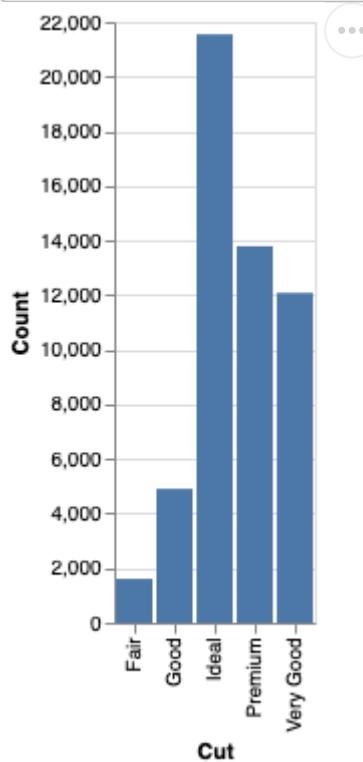
# Summarizing `cut` in a table

```
1 diamonds_cut = diamonds.groupby('cut').size()  
2 diamonds_cut
```

```
cut  
Fair          1610  
Good         4906  
Very Good    12082  
Premium      13791  
Ideal        21551  
dtype: int64
```

# Summarizing with a bar graph

```
1 diamonds_cut = diamonds_cut.reset_index().rename(columns={0:'N'})  
2  
3 alt.Chart(diamonds_cut).mark_bar().encode(  
4     alt.X('cut:O', title = "Cut"),  
5     alt.Y('N:Q', title = "Count")  
6 )
```



# Categorical variables: summary

- This section is very brief because there's basically only one good way to plot categorical variables with a small number of categories and this is it.
  - You can use `mark_point()` instead of `mark_bar()`, but overall, there's a clear right answer about how to do this.
- We include this material mainly to foreshadow the fact that we will do a lot on categorical variables in the next lecture when we get to “Exploring Co-variation”

# Continuous Variables

# Continuous variables: roadmap

- Binning + histograms using `movies`
- Histograms and density plots using `penguins`
- Exploring carat size using `diamonds`

Remark: The skills are absolutely fundamental and so we will intentionally be a bit repetitive.

# Recall: movies dataset

```
1 movies_url = 'https://cdn.jsdelivr.net/npm/vega-datasets@1/data/movies.json'  
2 movies = pd.read_json(movies_url)  
3 movies.head()
```

	Title	US_Gross	Worldwide_Gross	US_DVD_Sales	Production_Budget	Re
0	The Land Girls	146083.0	146083.0	NaN	8000000.0	Ju
1	First Love, Last Rites	10876.0	10876.0	NaN	300000.0	Au
2	I Married a Strange Person	203134.0	203134.0	NaN	250000.0	Au

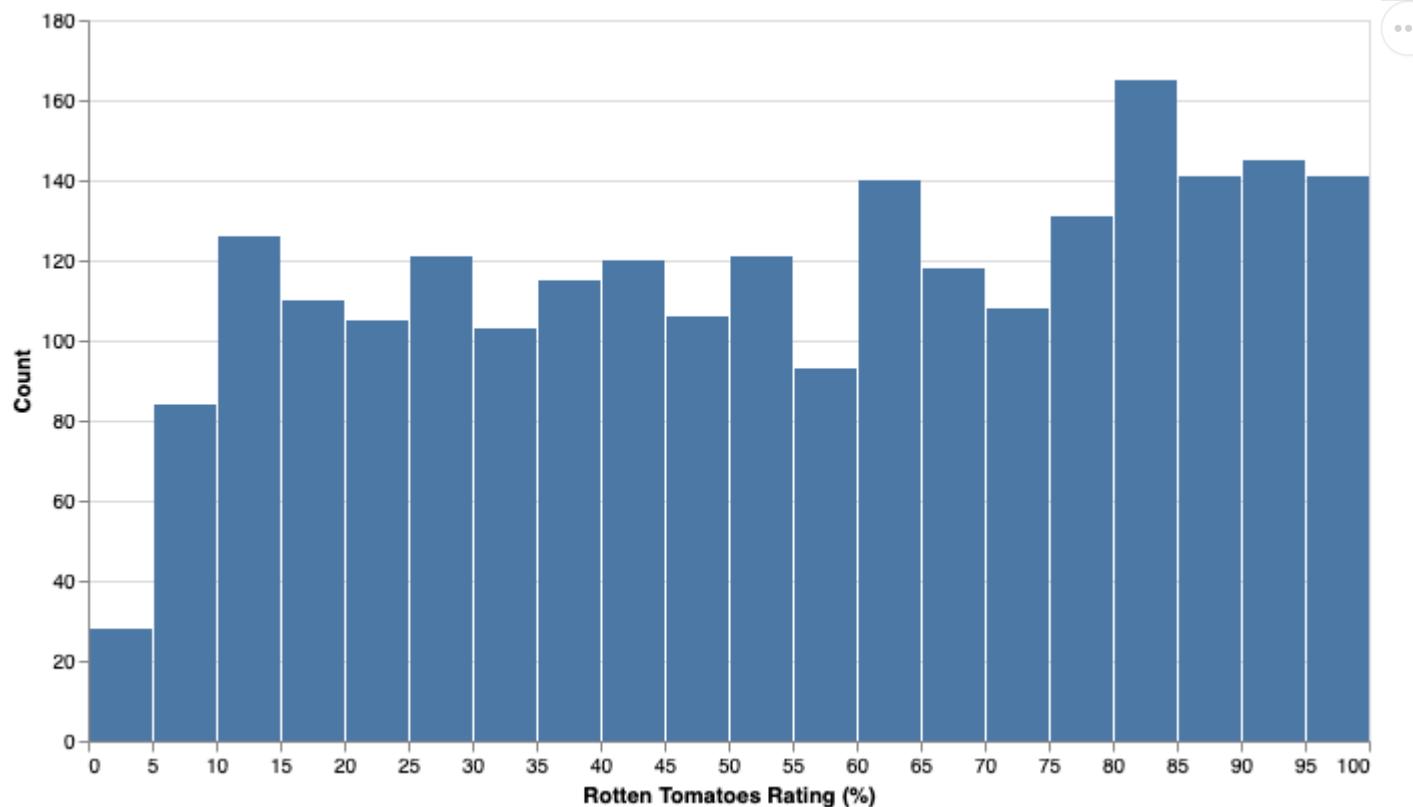
	Title	US_Gross	Worldwide_Gross	US_DVD_Sales	Production_Budget	Re
3	Let's Talk About Sex	373615.0	373615.0	NaN	300000.0	Se
4	Slam	1009819.0	1087521.0	NaN	1000000.0	Oc

# Histogram using `mark_bar()`

- Rotten Tomatoes ratings are determined by taking “thumbs up” and “thumbs down” judgments from film critics and calculating the percentage of positive reviews.
- This is a continuous measure, but we can bin it to create a **histogram** of frequencies

# Histogram using mark\_bar()

```
1 hist_rt = alt.Chart(movies_url).mark_bar().encode(  
2     alt.X('Rotten_Tomatoes_Rating:Q', bin=alt.BinParams(maxbins=20), title  
3     alt.Y('count():Q', title = "Count")  
4 )  
5 hist_rt
```

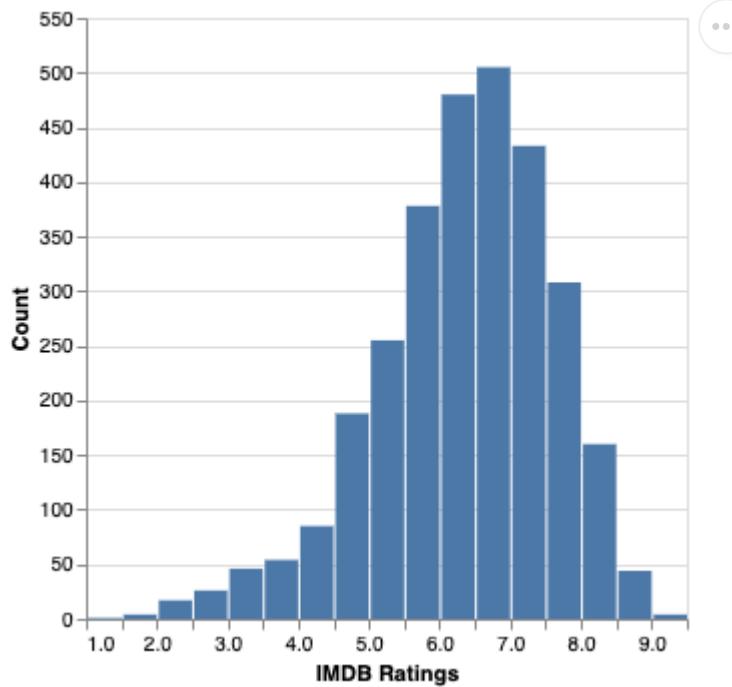


Discussion question: what are the headline and sub-messages?

# Histogram of IMDB ratings

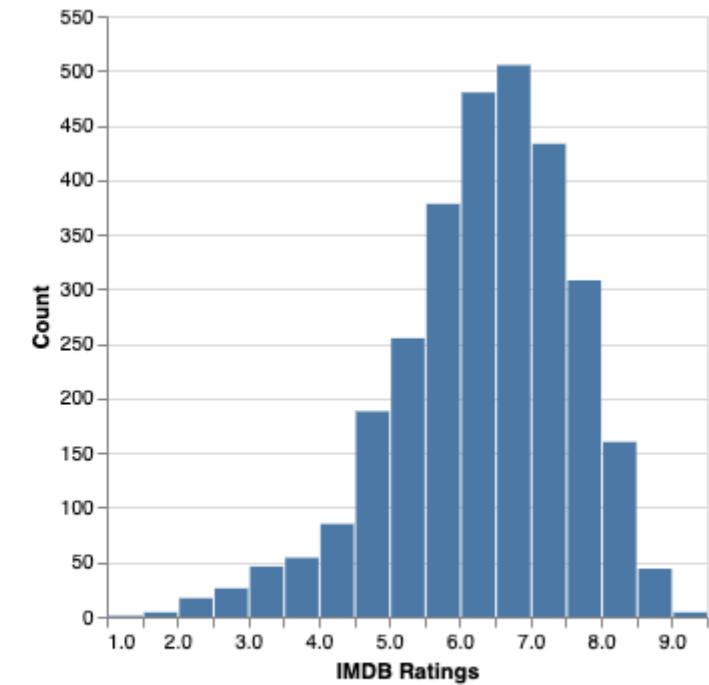
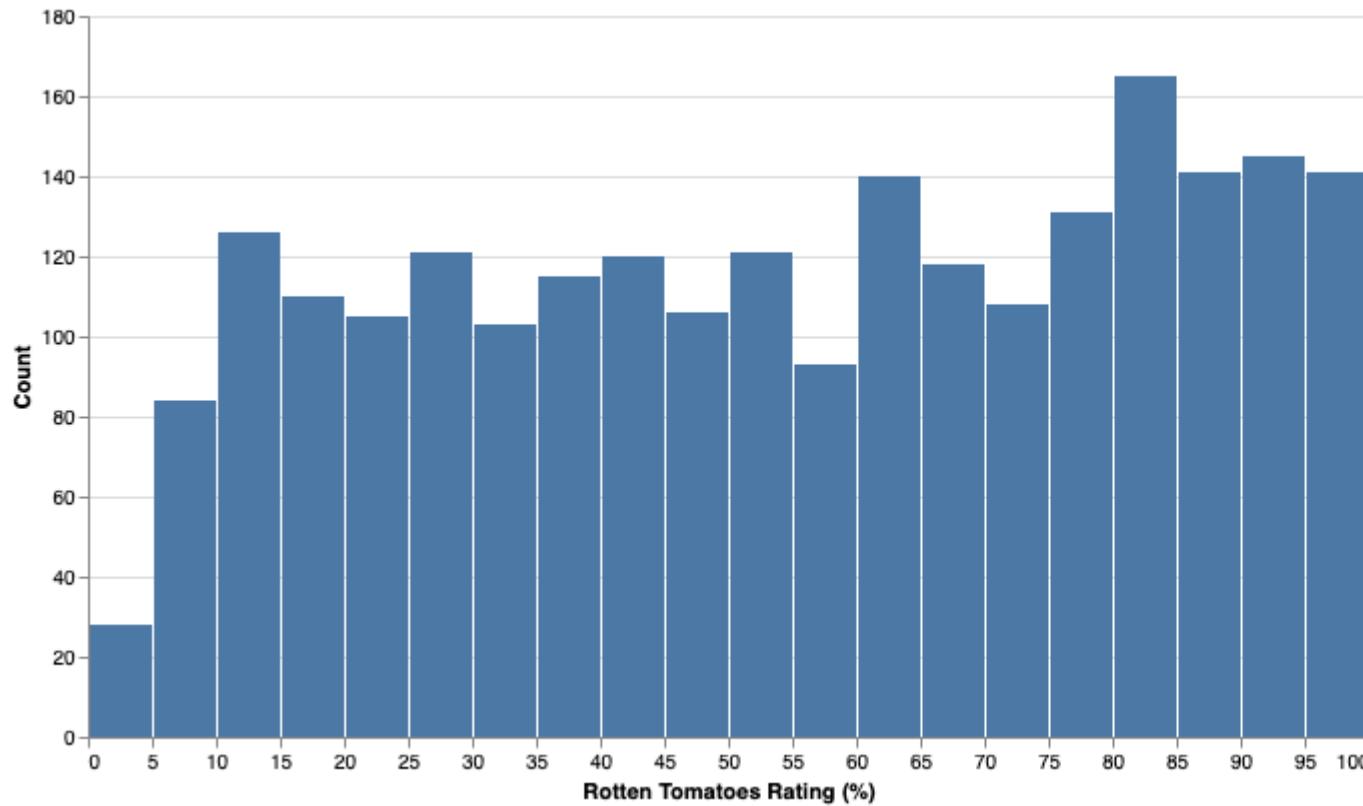
IMDB ratings are formed by averaging scores (ranging from 1 to 10) provided by the site's users.

```
1 hist_imdb = alt.Chart(movies_url).mark_bar().encode(
2     alt.X('IMDB_Rating:Q', bin=alt.BinParams(maxbins=20), title = "IMDB Rat
3     alt.Y('count():Q', title = "Count")
4 )
5 hist_imdb
```



# Side-by-side

```
1 hist_rt | hist_imdb
```



Discussion question: compare the two ratings distributions. If *your goal for the headline of the graph to be about differentiating between good and bad movies*, which is more informative?

# Introducing the `penguins` dataset

```
1 url = ("https://raw.githubusercontent.com/mcnakhaee/palmerpenguins/master/p
2 penguins = pd.read_csv(url)
3 penguins.head()
```

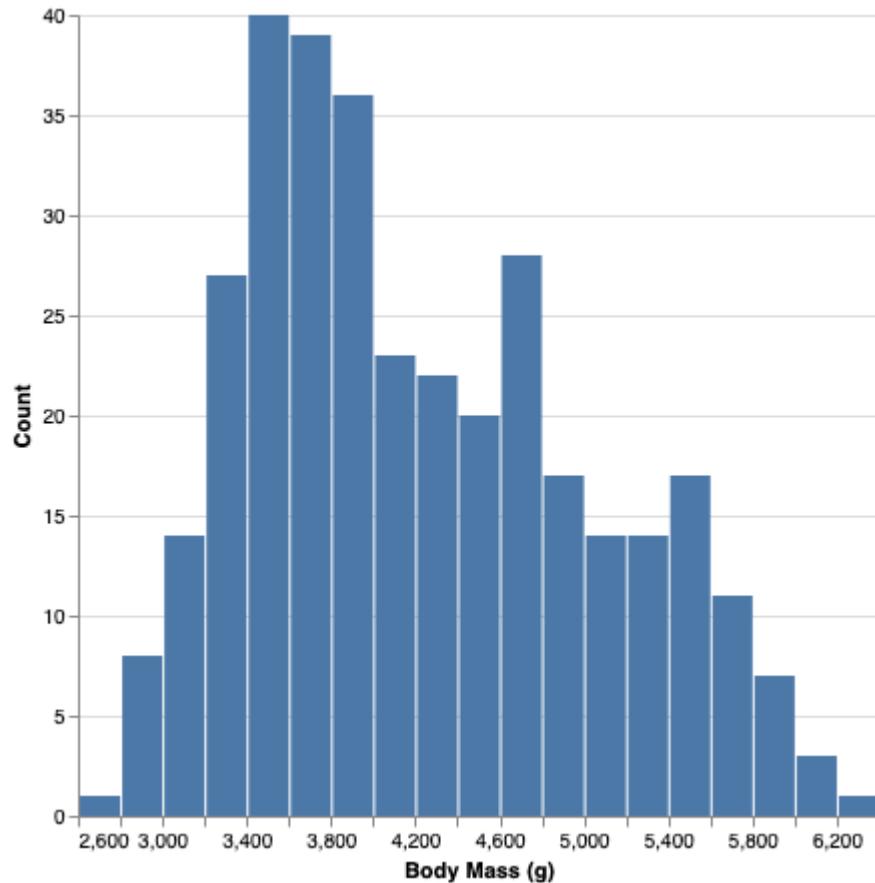
	species	island	bill_length_mm	bill_depth_mm	flipper_length_mm	boc
0	Adelie	Torgersen	39.1	18.7	181.0	375
1	Adelie	Torgersen	39.5	17.4	186.0	380
2	Adelie	Torgersen	40.3	18.0	195.0	325
3	Adelie	Torgersen	NaN	NaN	NaN	NaN
4	Adelie	Torgersen	36.7	19.3	193.0	345

# Histogram with steps of 200

- We previously picked the maximum number of equally-spaced bins (`BinParams(maxbins=20)`) and let `altair` choose “nice”-looking bin widths for the histogram
- Alternatively, we can manually control the bin width using `step`

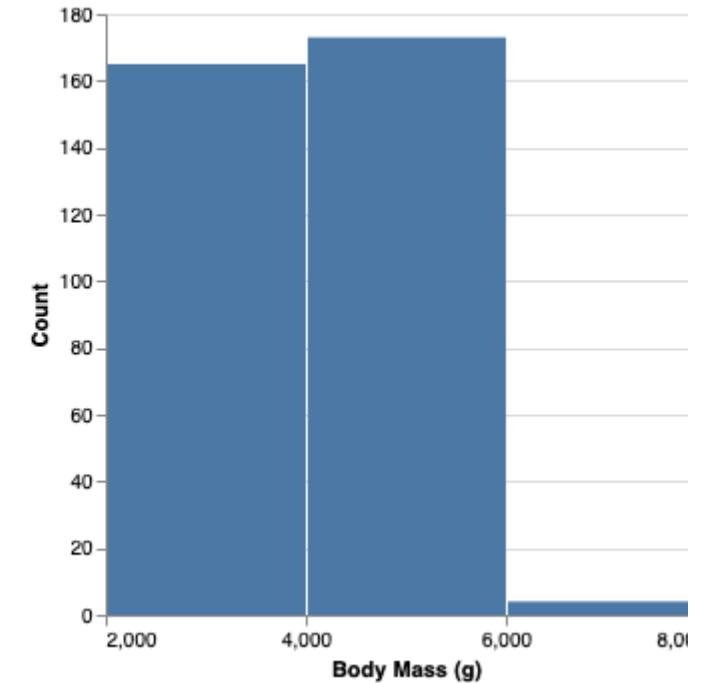
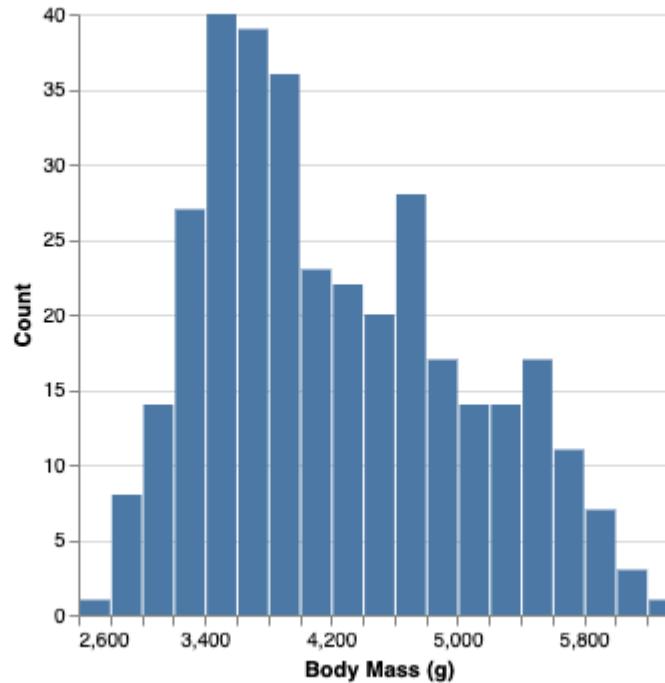
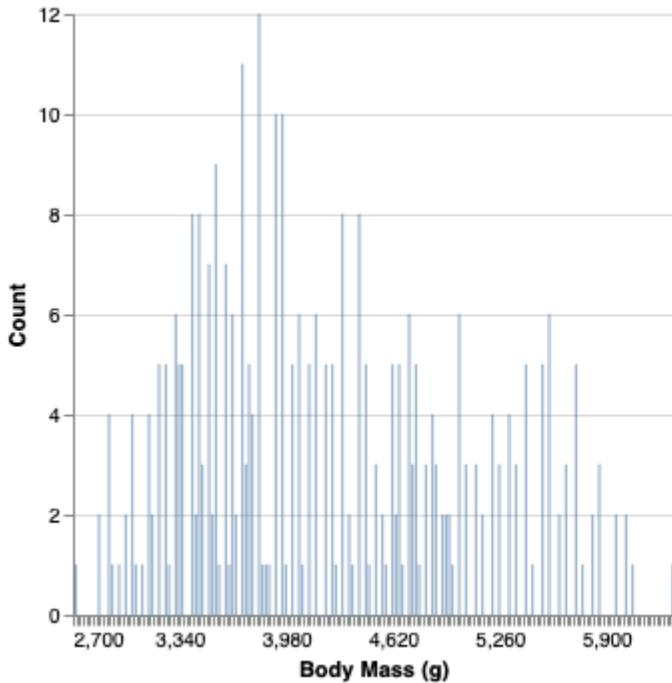
# Histogram with steps of 200

```
1 alt.Chart(penguins).mark_bar().encode(  
2     alt.X('body_mass_g:Q', bin=alt.BinParams(step=200), title = "Body Mass  
3     alt.Y('count()):Q', title = "Count")  
4 )
```



# Histogram step parameter

step=20 vs. step=200 vs. step=2000

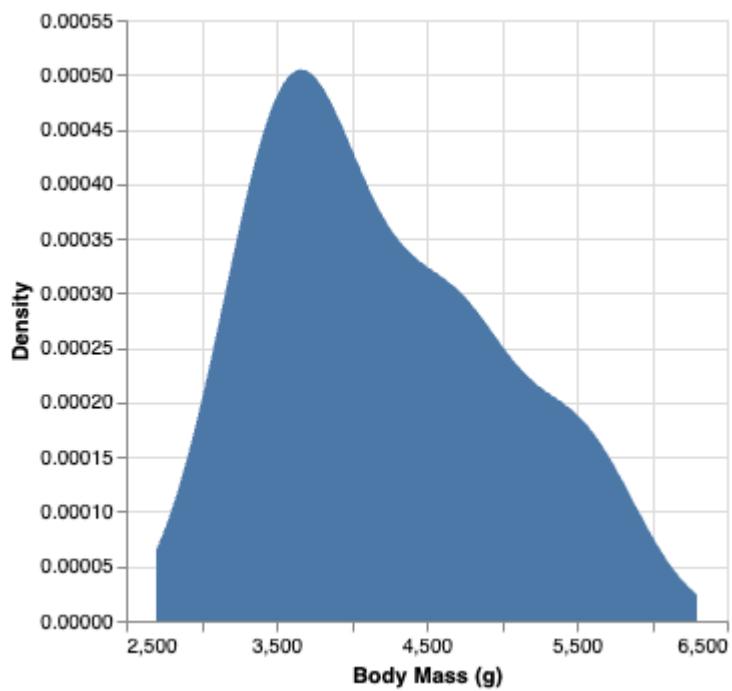


Discussion question: what message(s) come from each **binwidth** choice? Which do you prefer?

# Density plot

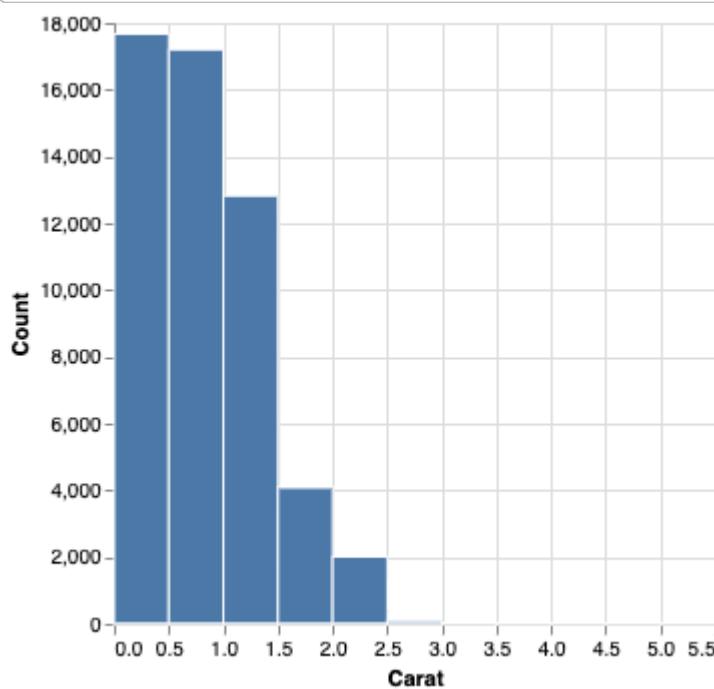
An alternative to a histogram for exploring frequency in continuous variable:  
density plot using `transform_density`

```
1 alt.Chart(penguins).transform_density(  
2     'body_mass_g',  
3     as_=['body_mass_g', 'density'])  
4 ).mark_area().encode(  
5     alt.X('body_mass_g:Q', title = "Body Mass (g)",  
6     alt.Y('density:Q', title = "Density")  
7 )
```



# Back to diamonds, focus on carat

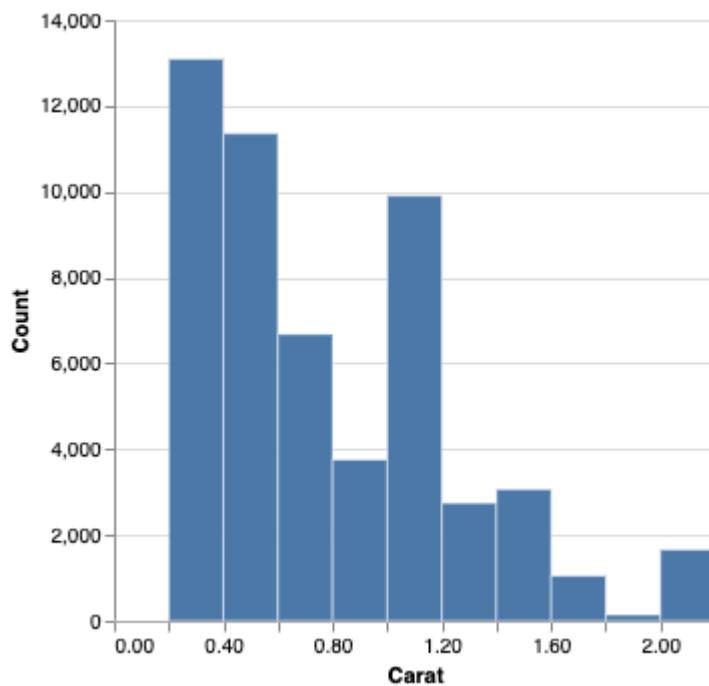
```
1 alt.data_transformers.disable_max_rows()  
2  
3 alt.Chart(diamonds).mark_bar().encode(  
4     alt.X('carat', bin=alt.Bin(maxbins=10), title = "Carat"),  
5     alt.Y('count()', title = "Count")  
6 )
```



First line disables `altair`'s maximum row limit (5,000)

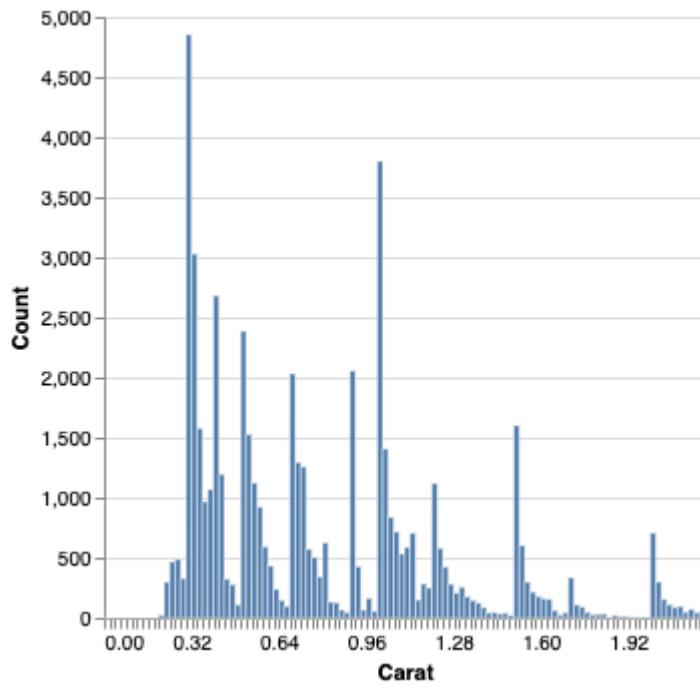
# Histogram of carat

```
1 diamonds_small = diamonds.loc[diamonds['carat'] < 2.1]
2
3 alt.Chart(diamonds_small).mark_bar().encode(
4     alt.X('carat', bin=alt.BinParams(step=0.2), title = "Carat"),
5     alt.Y('count()', title = "Count")
6 )
```



# In-class exercise: histogram of carat

```
1 alt.Chart(diamonds_small).mark_bar().encode(  
2     alt.X('carat', bin=alt.BinParams(step=0.02), title = "Carat"),  
3     alt.Y('count()', title = "Count"))
```



## Discussion questions

1. What is the headline of this plot? Submessages?
2. What questions does it raise?

# Typical continuous variables: summary

- Main tool to explore uni-dimensional continuous variables: histograms
- Varying the bin widths can reveal different patterns

# Continuous variables: unusual values

# Unusual continuous variables: roadmap

- case study: **y** dimension in diamonds
  - explore some unusual values
  - three options for handling unusual values

# diamonds: identify unusual y values

First pass to examine for unusual values: summary statistics

```
1 diamonds['y'].describe()
```

```
count      53940.000000
mean       5.734526
std        1.142135
min        0.000000
25%        4.720000
50%        5.710000
75%        6.540000
max        58.900000
Name: y, dtype: float64
```

# diamonds: examine unusual y values

```
1 diamonds.loc[(diamonds['y'] < 3) | (diamonds['y'] > 20)]
```

	carat	cut	color	clarity	depth	table	price	x	y	z	carat_bin	carat_bin_left	carat_k
11963	1.00	Very Good	H	VS2	63.3	53.0	5139	0.00	0.0	0.00	[1.0, 1.5)	1.0	1.5
15951	1.14	Fair	G	VS1	57.5	67.0	6381	0.00	0.0	0.00	[1.0, 1.5)	1.0	1.5
24067	2.00	Premium	H	SI2	58.9	57.0	12210	8.09	58.9	8.06	[2.0, 2.5)	2.0	2.5
24520	1.56	Ideal	G	VS2	62.2	54.0	12800	0.00	0.0	0.00	[1.5, 2.0)	1.5	2.0
26243	1.20	Premium	D	VVS1	62.1	59.0	15686	0.00	0.0	0.00	[1.0, 1.5)	1.0	1.5
27429	2.25	Premium	H	SI2	62.8	59.0	18034	0.00	0.0	0.00	[2.0, 2.5)	2.0	2.5
49189	0.51	Ideal	E	VS1	61.8	55.0	2075	5.15	31.8	5.12	[0.5, 1.0)	0.5	1.0
49556	0.71	Good	F	SI2	64.1	60.0	2130	0.00	0.0	0.00	[0.5, 1.0)	0.5	1.0
49557	0.71	Good	F	SI2	64.1	60.0	2130	0.00	0.0	0.00	[0.5, 1.0)	0.5	1.0

# iamonds: compare to 10 random diamonds

```
1 diamonds.sample(n=10)
```

	carat	cut	color	clarity	depth	table	price	x	y	z	carat_bin	carat_bin_left	carat_k
41108	0.59	Ideal	F	SI2	61.7	54.0	1195	5.38	5.44	3.34	[0.5, 1.0)	0.5	1.0
12327	1.20	Very Good	J	VS1	62.9	60.0	5226	6.64	6.69	4.19	[1.0, 1.5)	1.0	1.5
14318	1.32	Premium	J	VS2	60.4	58.0	5800	7.11	7.14	4.30	[1.0, 1.5)	1.0	1.5
32259	0.30	Ideal	E	VVS2	61.7	57.0	789	4.31	4.34	2.67	[0.0, 0.5)	0.0	0.5
1413	0.71	Premium	D	VS2	58.7	61.0	2968	5.88	5.85	3.44	[0.5, 1.0)	0.5	1.0
22496	1.71	Ideal	H	VS2	63.0	57.0	10534	7.57	7.53	4.76	[1.5, 2.0)	1.5	2.0
25764	2.03	Good	H	SI2	59.6	61.0	14720	8.16	8.21	4.88	[2.0, 2.5)	2.0	2.5
24677	1.52	Premium	G	VS1	61.4	58.0	13034	7.38	7.35	4.52	[1.5, 2.0)	1.5	2.0
13695	1.00	Very Good	G	VS2	61.8	59.0	5600	6.29	6.37	3.91	[1.0, 1.5)	1.0	1.5
24942	2.02	Fair	J	VS1	64.6	59.0	13399	7.95	7.90	5.12	[2.0, 2.5)	2.0	2.5

# What to do with unusual values?

1. Drop row
2. Code value to NA
3. Winsorize value

# Option 1: drop rows

```
1 diamonds_clean = diamonds.loc[(diamonds['y'] >= 3) | (diamonds['y'] <= 20)]  
2 diamonds_clean
```

	carat	cut	color	clarity	depth	table	price	x	y	z	carat_bin	carat_bin_left	carat_b
0	0.23	Ideal	E	SI2	61.5	55.0	326	3.95	3.98	2.43	[0.0, 0.5)	0.0	0.5
1	0.21	Premium	E	SI1	59.8	61.0	326	3.89	3.84	2.31	[0.0, 0.5)	0.0	0.5
2	0.23	Good	E	VS1	56.9	65.0	327	4.05	4.07	2.31	[0.0, 0.5)	0.0	0.5
3	0.29	Premium	I	VS2	62.4	58.0	334	4.20	4.23	2.63	[0.0, 0.5)	0.0	0.5
4	0.31	Good	J	SI2	63.3	58.0	335	4.34	4.35	2.75	[0.0, 0.5)	0.0	0.5
...	...	...	...	...	...	...	...	...	...	...	...	...	...
53935	0.72	Ideal	D	SI1	60.8	57.0	2757	5.75	5.76	3.50	[0.5, 1.0)	0.5	1.0
53936	0.72	Good	D	SI1	63.1	55.0	2757	5.69	5.75	3.61	[0.5, 1.0)	0.5	1.0
53937	0.70	Very Good	D	SI1	62.8	60.0	2757	5.66	5.68	3.56	[0.5, 1.0)	0.5	1.0
53938	0.86	Premium	H	SI2	61.0	58.0	2757	6.15	6.12	3.74	[0.5, 1.0)	0.5	1.0
53939	0.75	Ideal	D	SI2	62.2	55.0	2757	5.83	5.87	3.64	[0.5, 1.0)	0.5	1.0

53940 rows × 13 columns

# Option 2: recode to missing

```
1 diamonds_missing = diamonds.copy()
2 diamonds_missing['y'] = np.where((diamonds_missing['y'] < 3) | (diamonds_mi
3 diamonds_missing[diamonds_missing['y'].isna()])
```

	carat	cut	color	clarity	depth	table	price	x	y	z	carat_bin	carat_bin_left	carat_l
11963	1.00	Very Good	H	VS2	63.3	53.0	5139	0.00	NaN	0.00	[1.0, 1.5)	1.0	1.5
15951	1.14	Fair	G	VS1	57.5	67.0	6381	0.00	NaN	0.00	[1.0, 1.5)	1.0	1.5
24067	2.00	Premium	H	SI2	58.9	57.0	12210	8.09	NaN	8.06	[2.0, 2.5)	2.0	2.5
24520	1.56	Ideal	G	VS2	62.2	54.0	12800	0.00	NaN	0.00	[1.5, 2.0)	1.5	2.0
26243	1.20	Premium	D	VVS1	62.1	59.0	15686	0.00	NaN	0.00	[1.0, 1.5)	1.0	1.5
27429	2.25	Premium	H	SI2	62.8	59.0	18034	0.00	NaN	0.00	[2.0, 2.5)	2.0	2.5
49189	0.51	Ideal	E	VS1	61.8	55.0	2075	5.15	NaN	5.12	[0.5, 1.0)	0.5	1.0
49556	0.71	Good	F	SI2	64.1	60.0	2130	0.00	NaN	0.00	[0.5, 1.0)	0.5	1.0
49557	0.71	Good	F	SI2	64.1	60.0	2130	0.00	NaN	0.00	[0.5, 1.0)	0.5	1.0

# Option 3: winsorize

Winsorizing re-codes outliers to a numeric value, keeping them in the data.

To winsorize at 1 percent:

- Replace anything less than the 1st percentile with the 1st percentile
- Replace anything more than the 99th percentile with the 99th percentile

# Option 3: winsorize

```
1 diamonds_winsor = diamonds.copy()
2 pctile01 = diamonds_winsor['y'].quantile(0.01)
3 pctile99 = diamonds_winsor['y'].quantile(0.99)
4
5 print(f"1st Percentile: {pctile01}")
6 print(f"99th Percentile: {pctile99}")
```

```
1st Percentile: 4.04
99th Percentile: 8.34
```

# Option 3: winsorize

```
1 diamonds_winsor[ 'y_winsor' ] = np.where(diamonds_winsor[ 'y' ] < pctile01, pctile01,
2                                     np.where(diamonds_winsor[ 'y' ] > pctile99, pctile99, p
3 diamonds_winsor
```

	carat	cut	color	clarity	depth	table	price	x	y	z	carat_bin	carat_bin_left	carat_b
0	0.23	Ideal	E	SI2	61.5	55.0	326	3.95	3.98	2.43	[0.0, 0.5)	0.0	0.5
1	0.21	Premium	E	SI1	59.8	61.0	326	3.89	3.84	2.31	[0.0, 0.5)	0.0	0.5
2	0.23	Good	E	VS1	56.9	65.0	327	4.05	4.07	2.31	[0.0, 0.5)	0.0	0.5
3	0.29	Premium	I	VS2	62.4	58.0	334	4.20	4.23	2.63	[0.0, 0.5)	0.0	0.5
4	0.31	Good	J	SI2	63.3	58.0	335	4.34	4.35	2.75	[0.0, 0.5)	0.0	0.5
...	...	...	...	...	...	...	...	...	...	...	...	...	...
53935	0.72	Ideal	D	SI1	60.8	57.0	2757	5.75	5.76	3.50	[0.5, 1.0)	0.5	1.0
53936	0.72	Good	D	SI1	63.1	55.0	2757	5.69	5.75	3.61	[0.5, 1.0)	0.5	1.0
53937	0.70	Very Good	D	SI1	62.8	60.0	2757	5.66	5.68	3.56	[0.5, 1.0)	0.5	1.0
53938	0.86	Premium	H	SI2	61.0	58.0	2757	6.15	6.12	3.74	[0.5, 1.0)	0.5	1.0
53939	0.75	Ideal	D	SI2	62.2	55.0	2757	5.83	5.87	3.64	[0.5, 1.0)	0.5	1.0

53940 rows × 14 columns

# When might you winsorize?

An example from [Earnings Instability](#) paper by Ganong and coauthors.

The paper is trying to quantify how much earnings change from month to month for the typical US worker.

Consider the following fake data (next slide)

# Toy winsorization example

Suppose we have observations for earnings changes. 99% of the data follows a normal distribution with std. dev. 0.2 and 1% of the data is extremely large changes

last month (\$)	this month (\$)	% change	% change
600	600	0%	0%
600	570	-5%	5%
600	540	-10%	10%
600	630	5%	5%
...			
(99% of sample)			
...			
600	300	-50%	50%
6000	300	-95%	95%
300	600	100%	100%
300	6000	1900%	1900%

What is the standard deviation of the % change in earnings?

assumption	SD
do not winsorize	97.2%
winsorize at 50%	20.5%

Illustrative calculation [here](#)

When else is this useful? Income data, test scores, stock returns.

# Real-world winsorization example

Table A-2: Earnings Volatility Under Different Winsorization Choices

Specification	Lower bound	Upper bound	Std. dev.
Winsorize top/bottom 1% of all changes	-0.66	1.83	31%
Winsorize top/bottom 0.1% of nonzero changes	-0.93	13	60%
Winsorize top/bottom 0.5% of nonzero changes	-0.81	4	41%
Winsorize top/bottom 1% of nonzero changes	-0.71	2.39	34%
Winsorize top/bottom 5% of nonzero changes	-0.42	0.71	22%
Winsorize changes larger than 50%	Bottom 2% of data	Top 5% of data	20%

Notes: The variable is the percent change in pay.

Source: Table A-2 from [Earnings Instability](#) paper

# Pros + cons of each option

- Winsorizing
  - Manipulates the data values
  - Allows you to use that observation + that variable in your analysis
- Recoding to missing
  - Manipulates the data values
  - Allows you to use that observation – just not *that variable* – in your analysis
- Dropping the observation
  - Does not manipulate the data values
  - But can't use that observation *at all* in your analysis

# iamonds: what would you do?

- What would you do where  $x$ ,  $y$ , and  $z$  are all 0?
- What would you do where  $y > 20$ ?

# iamonds: what would we do?

There is often not a “right” answer or you won’t know the answer without talking to a data provider.

Our best guesses:

- Rows where `x`, `y`, and `z` are all zero: set to `NA`
- Rows where `y > 20`: `winsorize?` (hard to know for sure...)

# Unusual continuous values: summary

Problem	Action
Erroneous row	drop row
Erroneous cell	set to NA or winsorize

**How do I decide which problem I have?** Examine unusual values in context of other columns (same row) and other rows (same columns).

**How do I decide whether to set to NA or winsorize?** Ideally, ask your data provider what's going on with these values.

# Unusual values case study

# Unusual values case study: roadmap

- Introduce mpg dataset

- Research question 1

What is the relationship between engine size and gas mileage?"

- Research question 2

Why do some cars have better than typical mileage?

- Ad hoc identification of outliers

- Inspect fields describing outliers

- Uncover pattern

# Introducing the mpg dataset

- `manufacturer` – car maker (e.g., toyota, ford)
- `model` – specific model name
- `displ` – engine size (liters)
- `hwy` – gas mileage highway miles per gallon
- `class` – vehicle class (compact, suv, pickup, etc.)

# Introducing the `mpg` dataset

	manufacturer	model	displ	hwy	class
0	audi	a4	1.8	29	compact
1	audi	a4	1.8	29	compact
2	audi	a4	2.0	31	compact
3	audi	a4	2.0	30	compact
4	audi	a4	2.8	26	compact
...	...	...	...	...	...
229	volkswagen	passat	2.0	28	midsize
230	volkswagen	passat	2.0	29	midsize

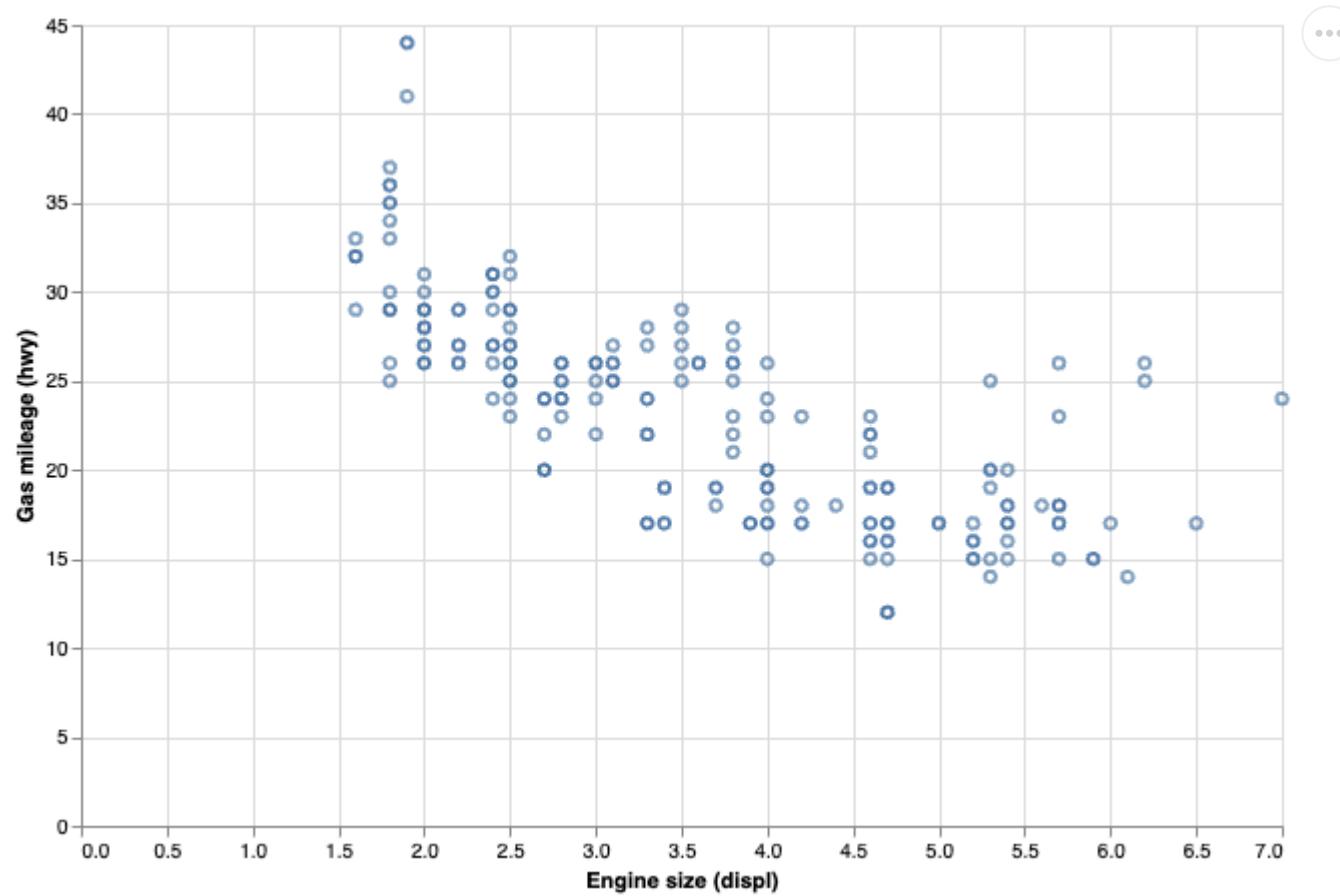
	manufacturer	model	displ	hwy	class
231	volkswagen	passat	2.8	26	midsize
232	volkswagen	passat	2.8	26	midsize
233	volkswagen	passat	3.6	26	midsize

234 rows × 5 columns

# Q1: What is the relationship between engine size and gas mileage?

```
1 base = alt.Chart(mpg).mark_point().encode(  
2     alt.X('displ:Q', title = "Engine size (displ)"),  
3     alt.Y('hwy:Q', title = "Gas mileage (hwy)")  
4 )  
5 base
```

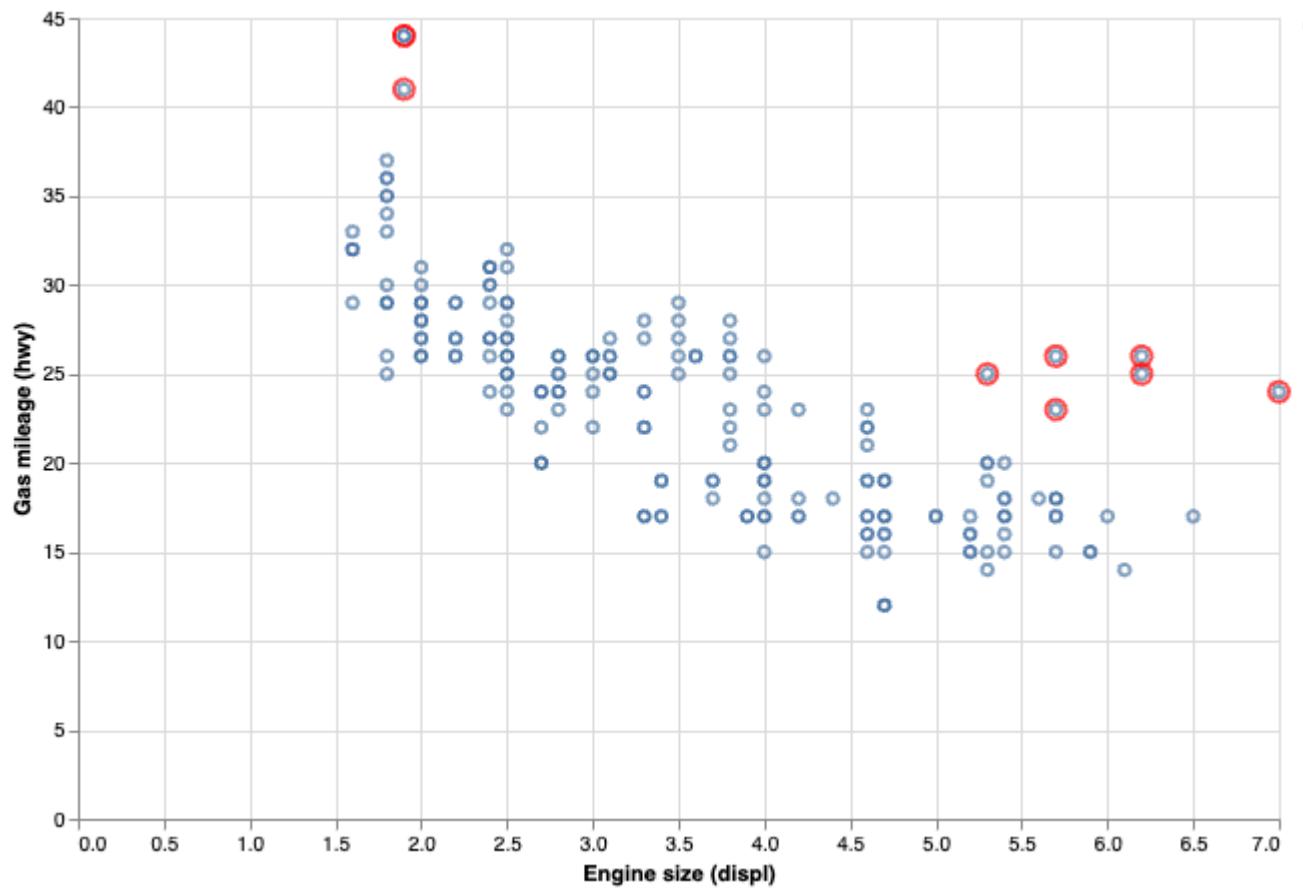
# Q1: What is the relationship between engine size and gas mileage?



# Q2: Why do some cars have better than typical mileage?

```
1 potential_outliers = mpg.loc[ (mpg["hwy"] > 40) | ((mpg["hwy"] > 20) & (mpg[
```

# Q2: Why do some cars have better than typical mileage? (plot)



# Q2: In-class exercise (table)

Discussion q – which fields do you want to study further on the plot (and why?)

```
[ 'manufacturer', 'model', 'displ', 'year', 'cyl', 'trans', 'drv', 'cty',
  'hwy', 'fl', 'class' ]
```

	manufacturer	model	displ	year	cyl	tr
23	chevrolet	corvette	5.7	1999	8	manual(l4)
24	chevrolet	corvette	5.7	1999	8	auto(l4)
25	chevrolet	corvette	6.2	2008	8	manual(l4)
26	chevrolet	corvette	6.2	2008	8	auto(s6)
27	chevrolet	corvette	7.0	2008	8	manual(l4)

	manufacturer	model	displ	year	cyl	tr
158	pontiac	grand prix	5.3	2008	8	auto(s4)
212	volkswagen	jetta	1.9	1999	4	manual(l)
221	volkswagen	new beetle	1.9	1999	4	manual(l)
222	volkswagen	new beetle	1.9	1999	4	auto(l4)

# Let's focus on two fields

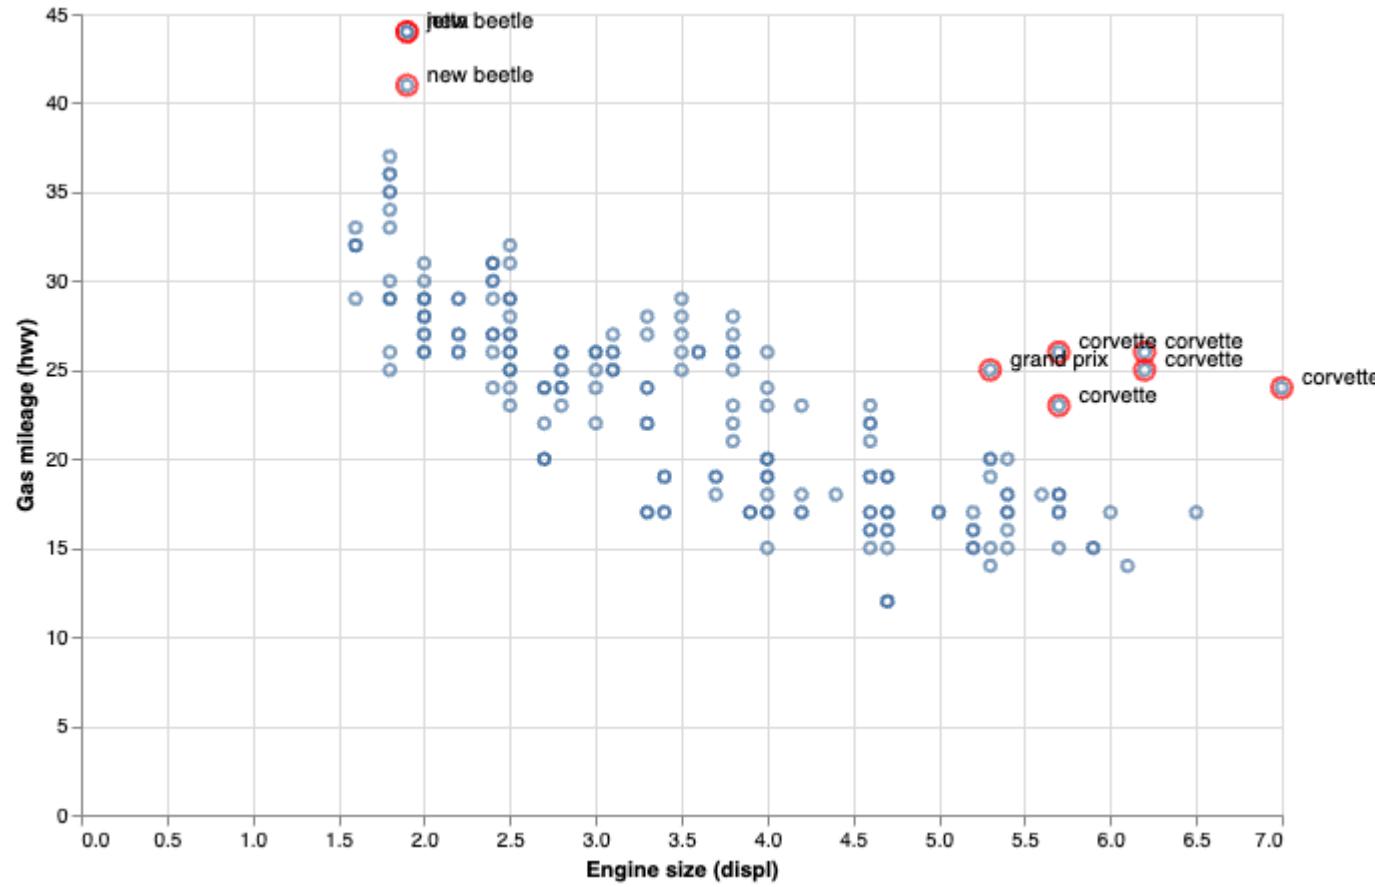
Fields

1. `model`
2. `class`

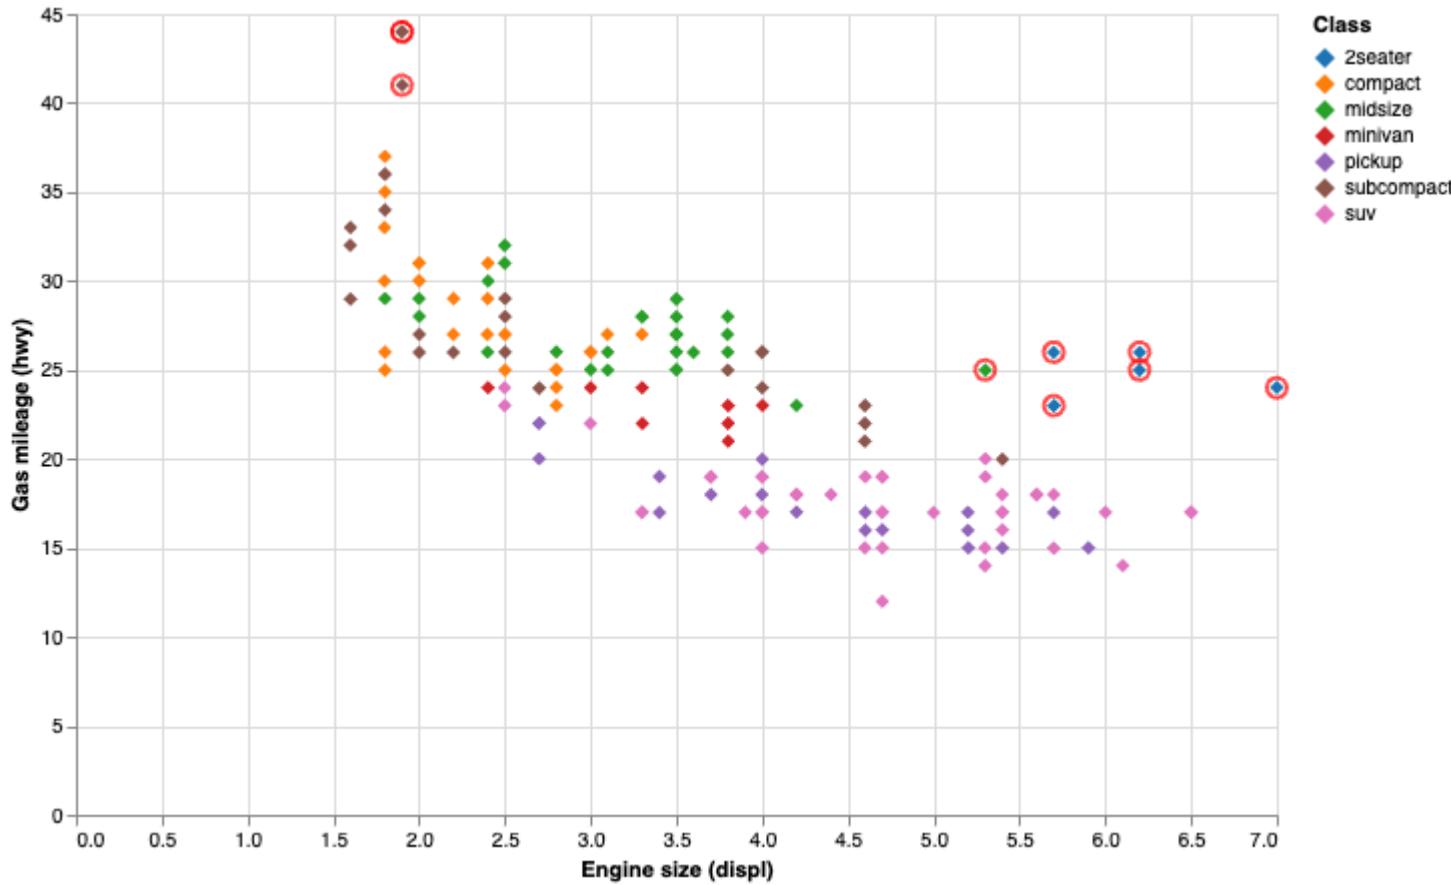
How did I know to use these? Context knowledge about different types of cars.

Don't have context knowledge about your dataset? Use LLM/Google/human subject matter expert to help you identify patterns

# Q2: Why do some cars have better than typical mileage?



# Q: How are there big engines and good mileage? color



# Discussion

- We applied labels to each outlier car using the `model` field
- We hypothesized that the field `class` would capture what these models would have in common
- This is the elegant corner case where one variable (`class`) explains many of the outlier patterns in terms of fuel-efficient cars
  - All the large engines on the top-right are `class == "2seater"`
  - Both of the small engines on the top-left are `class == "subcompact"` (but many subcompacts are less fuel efficient)
- Caveat: most datasets are not as clean as this example is, but we've chosen this example for instructional purposes

# Unusual values case study: summary

## Research question

Why do some cars have better than typical mileage? (What's going on with these outliers?)

## What did we do?

1. We identified outliers by hand
2. We looked at variables for those outliers
3. We went back to the plot with a variable which we thought could provide a unified explanation for the outliers.
4. It mostly did (subcompact cars and 2 seater cars are both very fuel efficient)